

Final Report for Replicated Concurrency Control and Recovery Project

Team Members:

- Shika Rao (sr7463)
- Khushboo . (kx2252)

Programming Language: Python

1. High-Level Architecture

The project simulates a replicated database system with snapshot isolation, available-copies replication, and first-committer-wins (FCW) conflict handling. The code is organized into three main modules:

- `data_structure.py` : domain data types (Version, VariableHistory, Site, Transaction).
- [`repcrec.py`](#) : core RepCRec engine implementing the concurrency-control and replication logic.
- [`runner.py`](#) : command-line driver: reads test scripts, creates a fresh engine per test, and invokes RepCRec methods.

```
runner.py
    ↓ (drives)
repcrec.py (RepCRec engine)
    ↓ uses
data_structure.py
    ↓ defines
{ Transaction, Site, VariableHistory, Version}
```

2. Component Overview

Data Structures used

`Version`

```
@dataclass
class Version:
```

```
value: int
commit_time: int
writer: Optional[str]
sites: Set[int]
```

- Represents a committed version of a variable.
- `commit_time` is a global logical time used to enforce snapshot isolation.
- `sites` lists the sites that have this version (for replicated variables).

VariableHistory

```
@dataclass
class VariableHistory:
    name: str
    versions: List[Version] = field(default_factory=list)

    def latest_before(self, ts: int) -> Optional[Version]:
        ...
```

- Maintains the history of committed versions of a variable ($x_1 \dots x_{20}$).
- `latest_before(ts)` is the core snapshot helper: it picks the most recent committed version $\leq ts$, and is used during reads to implement snapshot isolation.

Site

```
@dataclass
class Site:
    id: int
    is_up: bool = True
    data: Dict[str, int] = field(default_factory=dict)
    can_read: Dict[str, bool] = field(default_factory=dict)
    failure_times: List[int] = field(default_factory=list)
    recovery_times: List[int] = field(default_factory=list)

    def fail(self, time: int) -> None: ...
    def recover(self, time: int, is_replicated) -> None: ...
```

- Represent a physical site (1–10).

- Track whether the site is up or down.
- Track committed data visible at that site in `data[var]`.
- Track readability of replicated variables after recovery via `can_read[var]`:
 - For replicated vars, `can_read[var]` is `False` after recovery until a new commit touches that var at the site.
 - For non-replicated vars, `can_read[var]` is `True` immediately after recovery.
- Maintain `failure_times` and `recovery_times`, which the engine uses to check “site up continuously between time A and time B.”

Transaction

```
@dataclass
class Transaction:
    tid: str
    start_time: int
    status: str = "active" # active | committed | aborted | waiting
    (conceptually)
    read_vars: Set[str] = field(default_factory=set)
    write_buffer: Dict[str, int] = field(default_factory=dict)
    write_sites: Dict[str, Set[int]] = field(default_factory=dict)
    site_write_times: Dict[int, int] = field(default_factory=dict)
    commit_time: Optional[int] = None
```

Responsibilities:

- Represent a logical transaction T_i :
 - `start_time` is the snapshot time.
 - `write_buffer` holds uncommitted writes.
 - `write_sites[var]` stores which sites should receive that variable at commit.
 - `site_write_times` tracks the earliest write time to each site, needed for the available-copies failure rule.
- `status` is updated by the engine as commands execute.

Replicated Concurrency Control Engine

```
class RepCRec:
    def __init__(self):
        self.time = 0
        self.sites: Dict[int, Site] = {}
```

```

    self.vars: Dict[str, VariableHistory] = {}
    self.txns: Dict[str, Transaction] = {}
    self.var_last_writer: Dict[str, Tuple[str, int]] = {}
    self._init_sites_and_vars()

```

Global State

- time: global logical clock incremented once per input command line.
- sites: all Site objects.
- vars: all VariableHistory objects for x1...x20.
- txns: active/finished Transaction objects.
- var_last_writer[var] = (tid, commit_time): used for **first-committer-wins**.

`begin(tid)`

- execute_line() parses a line of the test script and routes it
- Advances self.time by 1 for each real command.
Creates a new Transaction with current self.time as start_time.
- Stores it in self.txns.

`read(tid, var)`

Implements snapshot isolation and available copies:

1. Fetch the transaction; ignore reads for aborted/committed transactions.
2. Use VariableHistory.latest_before(txn.start_time) to get the snapshot version.
3. If variable is non-replicated:
 - Read from its home site if it is up and the snapshot version exists there.
4. If variable is replicated:
 - Try to find a site s such that:
s is up, var is readable at s (can_read[var] is True if replicated), s is in snapshot.sites, and site_up_continuously(s, snapshot.commit_time, txn.start_time) holds.
 - If none exists:
Either abort immediately (for some cases), or conceptually mark the transaction as “waiting for a snapshot”.

`write(tid, var, value)`

1. Determine sites that currently hold the variable and are up:
 - o Replicated: all up sites.
 - o Non-replicated: only home site, if up.
2. If no such sites: transaction is aborted (project spec says it should “wait”; we implement an eager abort as a simplification).
3. Buffer the write in `txn.write_buffer[var] = value`.
4. Record `txn.write_sites[var]` and update `txn.site_write_times[site]` for available-copies failure rules.
5. Print that the write is buffered (no visible state change yet).

`end(tid)`

1. Available copies abort rule:

For each site s that this transaction wrote to:

 - o If s failed at any time f with $\text{write_time} < f \leq \text{current_time}$, abort the transaction.
2. First-committer-wins (FCW):

For each variable var in `write_buffer`:

 - o If $\text{var_last_writer}[\text{var}] = (\text{other_tid}, \text{other_ctime})$ with $\text{other_ctime} > \text{txn.start_time}$ and $\text{other_tid} \neq \text{tid}$, abort.
3. SSI / cycle detection: RW-cycle detection component:
 - o The engine tracks read-write relationships between transactions and aborts a transaction if it would close a dangerous cycle. This is implemented conservatively to avoid over-aborting: FCW + snapshot isolation already enforces most constraints.
4. If none of the above cause an abort, commit:
 - o Set `txn.commit_time = current_time`.
 - o For each $(\text{var}, \text{value})$ in `write_buffer`:
 - Apply the update to each site in `txn.write_sites[var]` that is up.
 - For replicated variables, mark `site.can_read[var] = True` on those sites.
 - Append a new Version to `vars[var].versions`.
 - Update `var_last_writer[var] = (tid, commit_time)`.
 - o Set `txn.status = "committed"` and print “Ti commits”.

Main Runner file

- Input test file with all the test cases. (test.txt contains all 25 sample test cases)
- Command: python3 [runner.py](#) test.txt

3. Reprozip replay in new directory:

```
repronzip directory setup repcrec.rpz repro_test_run  
repronzip directory run repro_test_run
```